



1




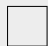








Profs. Nicolas Flammarion and Martin Jaggi
Machine Learning – CS-433 - IC
20.01.2022 from 08h15 to 11h15 in STCC
Duration : 180 minutes

Student One

SCIPER: 111111

Do not turn the page before the start of the exam. This document is double-sided, has 20 pages, the last ones are possibly blank. Do not unstaple.

- This is a closed book exam. No electronic devices of any kind.
- Place on your desk: your student ID, writing utensils, one double-sided A4 page cheat sheet (hand-written or 11pt min font size) if you have one; place all other personal items below your desk.
- You each have a different exam.
- Only answers in this booklet count. No extra loose answer sheets. You can use the last two pages as scrap paper.
- For the **multiple choice** questions, we give :
 - +2 points if your answer is correct,
 - 0 points if you give no answer or more than one,
 - −0.5 points if your answer is incorrect.
- For the **true/false** questions, we give :
 - +1 points if your answer is correct,
 - 0 points if you give no answer or more than one,
 - −1 points if your answer is incorrect.
- Use a **black or dark blue ballpen** and clearly erase with **correction fluid** if necessary.
- If a question turns out to be wrong or ambiguous, we may decide to nullify it.

Respectez les consignes suivantes Observe this guidelines Beachten Sie bitte die unten stehenden Richtlinien		
choisir une réponse select an answer Antwort auswählen	ne PAS choisir une réponse NOT select an answer NICHT Antwort auswählen	Corriger une réponse Correct an answer Antwort korrigieren
  		 
ce qu'il ne faut PAS faire what should NOT be done was man NICHT tun sollte		
     		



First part: multiple choice questions

For each question, mark the box corresponding to the correct answer. Each question has **exactly one** correct answer.

Robustness to outliers

We consider a classification problem on linearly separable data. Our dataset had an outlier—a point that is very far from the other datapoints in distance (and also far from margins in SVM but still correctly classified by the SVM classifier).

We trained the SVM, logistic regression and 1-nearest-neighbour models on this dataset. We tested trained models on a test set that comes from the same distribution as training set, but doesn't have any outlier points. After that we removed the outlier and retrained our models.

Question 1 After retraining, which classifier will **change** its decision boundary around the test points.

- ☒ Logistic regression
- ☐ SVM
- ☐ 1-nearest-neighbors classifier
- ☐ All of them

Solution: The solution of the SVM problem only depends on the support vectors, and outlier is not a support vector as stated in the problem. Thus SVM decision boundary would not change after retraining. For 1-NN classifier decisions are based on the nearest datapoints that would be not outlier points, as outlier is very far from other datapoints. Thus decisions of 1-NN classifier would also not change. Logistic regression solution is based on all of the datapoints thus removing one point will change the decision boundary.

SVM

Question 2 For any vector $\mathbf{v} \in \mathbb{R}^D$ let $\|\mathbf{v}\|_2 := \sqrt{v_1^2 + \dots + v_D^2}$ denote the Euclidean norm. The hard-margin SVM problem for linearly separable points in \mathbb{R}^D is to minimize the Euclidean norm $\|\mathbf{w}\|_2$ under some constraints. What are the additional constraints for this optimization problem?

- ☐ $\mathbf{w}^\top \mathbf{x}_n \geq 1 \quad \forall n \in \{1, \dots, N\}$
- ☐ $\frac{y_n}{\mathbf{w}^\top \mathbf{x}_n} \geq 1 \quad \forall n \in \{1, \dots, N\}$
- ☒ $y_n \mathbf{w}^\top \mathbf{x}_n \geq 1 \quad \forall n \in \{1, \dots, N\}$
- ☐ $y_n + \mathbf{w}^\top \mathbf{x}_n \geq 1 \quad \forall n \in \{1, \dots, N\}$



Cross Validation

Question 3 Consider the K -fold cross validation on a linear regression model with a sufficiently large amount of training data. When K is large, the computational complexity of the K -fold cross validation with respect to K is of order

- ☐ $\mathcal{O}(1/K)$.
- ☒ $\mathcal{O}(K)$.
- ☐ $\mathcal{O}(1)$.
- ☐ $\mathcal{O}(K(K-1))$.

Solution: Because each training process uses $(K-1)/K = O(1)$ fraction of the data, and there are K such processes.

Bias-Variance Decomposition

Question 4 Consider a regression model where data (x, y) is generated by input x uniformly randomly sampled from $[0, 1]$ and $y(x) = x^2 + \varepsilon$, where ε is random noise with mean 0 and variance 1. Two models are carried out for regression: model A is a trained quadratic function $g(x; \mathbf{w}) = w_2x^2 + w_1x + w_0$ where $\mathbf{w} = (w_0, w_1, w_2)^\top \in \mathbb{R}^3$, and model B is a constant function $h(x) = 1/2$. Then compared to model B, model A has

- ☐ higher bias, higher variance.
- ☒ lower bias, higher variance.
- ☐ higher bias, lower variance.
- ☐ lower bias, lower variance.

Solution: Model B has zero variance because it outputs a constant $1/2$ which is not related to the training data.

Optimization

Question 5 Let n be an integer such that $n \geq 2$ and let $A \in \mathbb{R}^{n \times n}$, and $\mathbf{x} \in \mathbb{R}^n$, consider the function $f(\mathbf{x}) = \mathbf{x}^\top A \mathbf{x}$ defined over \mathbb{R}^n . Which of the following is the gradient of the function f ?

- ☒ $A^\top \mathbf{x} + A \mathbf{x}$
- ☐ $2A^\top \mathbf{x}$
- ☐ $2\mathbf{x}^\top A$
- ☐ $2A \mathbf{x}$

Solution: $\nabla f(\mathbf{x}) = A^\top \mathbf{x} + A \mathbf{x}$. See lab 11 exercise 1, here the matrix A is not symmetric.



Question 6 Which of the following functions reaches a global maximum on the set I ? (Note that $[\cdot, \cdot]$ and (\cdot, \cdot) denote closed and open intervals respectively)

- $f_1(x) = -x^4$, $I = [-5, 5]$
- $f_2(x) = \arccos(x)$, $I = (-1, 1)$
- $f_3(x) = x \exp(-x)$, $I = (-\infty, 0)$
- $f_4(x) = \sin(\cos(x)) \sin(x)$, $I = \mathbb{R}_+$

☒ f_1, f_4

☐ f_1, f_3, f_4

☐ f_1, f_2, f_3, f_4

☐ f_1, f_2, f_4

Solution:

- f_1 : Yes, computing the first and second derivative leads to the maximum being reached at $x = 0$. Also note that any continuous function on compact set reaches its maximum and minimum.
- f_2 : No, we can compute the first derivative starting from:

$$\begin{aligned}\cos(\arccos(x)) &= x \\ -\sin(\arccos(x)) \arccos'(x) &= 1 \\ \arccos'(x) &= -\frac{1}{\sin(\arccos(x))} \\ \arccos'(x) &= -\frac{1}{\sqrt{1 - \cos(\arccos(x))^2}} \\ \arccos'(x) &= -\frac{1}{\sqrt{1 - x^2}}\end{aligned}$$

The derivative of f can not be equal to 0 on I .

- f_3 : No, the first derivative is $-e^x(x - 1)$, hence no maximum can be reached on I .
- f_4 : Yes, since f is the product of periodic functions ($\sin(\cos(x))$ and $\sin(x)$), it is also periodic. Since f is periodic on I of not finite measure, f reaches its maximum infinitely many times.
- f_5 : Yes, Computing the gradient of f , we find $\nabla f(x_1, x_2) = (-\operatorname{sgn}(x_1) - 6x_2, -10x_2)$, which is equal to $(0, 0)$ for $(x_1, x_2) = (0, 0)$. Then computing the Hessian, we find that it is a diagonal matrix filled with negative values, hence f does reach a maximum at $(0, 0)$.



Linear Models

Question 7 Consider a linear regression model on a dataset which we split into a training set and a test set. After training, our model gives a mean-squared error of 0.1 on the training set and a mean-squared error of 5.3 on the test set. Recall that the mean-squared error (MSE) is given by:

$$MSE_{\mathbf{w}}(\mathbf{y}, \mathbf{X}) = \frac{1}{2N} \sum_{n=1}^N (y_n - \mathbf{x}_n^{\top} \mathbf{w})^2$$

Which of the following statements is **correct** ?

- ☒ Ridge regression can help reduce the gap between the training MSE and the test MSE.
- ☐ Retraining the model with feature augmentation (e.g. adding polynomial features) will increase the training MSE.
- ☐ Retraining while discarding some training samples will likely reduce the gap between the train MSE and the test MSE.
- ☐ Using cross-validation can help decrease the training MSE of this very model.

Solution:

- *feature augmentation*: **Incorrect**, using feature augmentation will increase overfitting, hence decrease the training MSE even more.
- *cross-validation*: **Incorrect**, cross-validation can help to select a model that overfits less, it does not help to get better performance on a specific model.
- *discarding some training samples*: **Incorrect**, reducing the number of training samples is more likely to increase the overfitting.
- *Ridge regression*: **Correct**, regularization from ridge regression can be useful to reduce overfitting.

Regularization

Question 8 Which of the following statements is **incorrect** ?

Training a model with L_1 -regularization ...

- ☐ can reduce the storage cost of the final model.
- ☒ is used to help escaping local minima during training.
- ☐ can reduce overfitting.
- ☐ can be named Lasso regression when in combination with an MSE loss function and a linear model.

Solution:

- *reduce the storage*: is **true**, the parameters tend to be sparse when using L_1 -regularization, thus parameters with 0 value do not need to be stored.
- *escaping local minima*: is **false** (so check this box), L_1 -regularization has nothing to do with the optimization of the model.
- *reduce overfitting*: is **true**, it is a regularization technique, hence it reduces the complexity of the model.
- *Lasso*: is **true**, see the course.



Neural networks

Let $f : \mathbb{R}^D \rightarrow \mathbb{R}$ be an L -hidden layer multi-layer perceptron (MLP) such that

$$f(\mathbf{x}) = \sigma_{L+1}(\mathbf{w}^\top \sigma_L(\mathbf{W}_L \sigma_{L-1}(\mathbf{W}_{L-1} \dots \sigma_1(\mathbf{W}_1 \mathbf{x}))),$$

with $\mathbf{w} \in \mathbb{R}^M$, $\mathbf{W}_1 \in \mathbb{R}^{M \times D}$ and $\mathbf{W}_\ell \in \mathbb{R}^{M \times M}$ for $\ell = 2, \dots, L$, and σ_i for $i = 1, \dots, L+1$ is an entry-wise activation function. For any MLP f and a classification threshold τ let $C_{f,\tau}$ be a binary classifier that outputs YES for a given input \mathbf{x} if $f(\mathbf{x}) \leq \tau$ and NO otherwise.

Question 9 Assume σ_{L+1} is the element-wise **sigmoid** function and $C_{f,\frac{1}{2}}$ is able to obtain a high accuracy on a given binary classification task T . Let g be the MLP obtained by multiplying the parameters **in the last layer** of f , i.e. \mathbf{w} , by 2. Moreover, let h be the MLP obtained by replacing σ_{L+1} with element-wise **ReLU**. Finally, let q be the MLP obtained by doing both of these actions. Which of the following is true?

$$\text{ReLU}(x) = \max\{x, 0\}$$

$$\text{Sigmoid}(x) = \frac{1}{1 + e^{-x}}$$

- ☐ $C_{h,0}$ may have an accuracy significantly lower than $C_{f,\frac{1}{2}}$ on T
- ☒ $C_{g,\frac{1}{2}}$, $C_{h,0}$, and $C_{q,0}$ have the same accuracy as $C_{f,\frac{1}{2}}$ on T
- ☐ $C_{g,\frac{1}{2}}$ may have an accuracy significantly lower than $C_{f,\frac{1}{2}}$ on T
- ☐ $C_{q,0}$ may have an accuracy significantly lower than $C_{f,\frac{1}{2}}$ on T

Solution: Since the threshold $\frac{1}{2}$ for sigmoid corresponds to the input to the last activation function being positive, $C_{h,0}$ is true. Moreover, multiplying the weights by 2 does not change the sign of the output. Therefore both $C_{g,\frac{1}{2}}$ and $C_{q,0}$ are also true.

Question 10 Assume the weights in f were obtained by initializing **all parameters to zero** and running SGD. Assume there exists τ such that $C_{f,\tau}$ is a good classifier. Let g be the MLP obtained by randomly keeping one neuron per layer and removing the other $M - 1$ neurons from each layer, and multiplying all weights except for the first layer by M (the number of neurons in each layer in the original network). What is the probability that $C_{g,\tau}$ is a good classifier?

- ☐ Less than $\frac{1}{M^L}$
- ☐ Between $\frac{1}{M^L}$ and $\frac{1}{M}$
- ☒ 1
- ☐ Not possible to determine with the given information.

Solution: When all weights are initialized to zero, the value for all hidden neurons will be the same and remain the same. As such, keeping one neuron in each layer and multiplying its output by the number of neurons will keep the output unchanged. Therefore g and f are equal.



Question 11 Which of the following techniques do *not* improve the generalization performance in deep learning?

- ☒ None. All techniques here improve generalization.
- ☐ Dropout
- ☐ Data augmentation
- ☐ Tuning the optimizer
- ☐ L2 regularization

Solution: Once correct hyperparameters are chosen, any of these strategies may be applied to improve the generalization performance.

Adversarial Examples

Consider a linear model $\hat{y} = \mathbf{x}^\top \mathbf{w}$ with the squared loss under an ℓ_∞ -bounded adversarial perturbation. For a single point (\mathbf{x}, y) , it corresponds to the following objective:

$$\max_{\tilde{\mathbf{x}}: \|\mathbf{x} - \tilde{\mathbf{x}}\|_\infty \leq \varepsilon} (y - \tilde{\mathbf{x}}^\top \mathbf{w})^2, \quad (\text{OP})$$

where $\|\mathbf{x} - \tilde{\mathbf{x}}\|_\infty \leq \varepsilon$ denotes the ℓ_∞ -norm, i.e. $|x_i - \tilde{x}_i| \leq \varepsilon$ for every i .

Question 12 Assume that $\mathbf{w} = (3, -2)^\top$, $\mathbf{x} = (-1, 2)^\top$, $y = 2$. What is the maximum value of the optimization problem in Eq. (OP)?

- ☒ $(9 + 5\varepsilon)^2$
- ☐ $(3 + 10\varepsilon)^2$
- ☐ $(10 - \varepsilon)^2$
- ☐ Other
- ☐ $(5 + 9\varepsilon)^2$

Solution: First, it's convenient to reparametrize the objective in terms of an additive perturbation $\boldsymbol{\delta}$: $\max_{\boldsymbol{\delta}: \|\boldsymbol{\delta}\|_\infty \leq \varepsilon} (y - \mathbf{x}^\top \mathbf{w} - \boldsymbol{\delta}^\top \mathbf{w})^2$. If we plug the given values of \mathbf{w} , \mathbf{x} , y , we get: $\max_{\boldsymbol{\delta}: \|\boldsymbol{\delta}\|_\infty \leq \varepsilon} (9 - 3\delta_1 + 2\delta_2)^2$. We can maximize this objective independently over δ_1 and δ_2 by noting that the optimal value is attained at a boundary of the feasible set, i.e. for $|\delta_1| = |\delta_2| = \varepsilon$. This leads to the maximizer $\boldsymbol{\delta}^* = (-\varepsilon, \varepsilon)^\top$ and the maximum value $(9 + 5\varepsilon)^2$.

Question 13 Assume that $\mathbf{w} = (3, -2)^\top$, $\mathbf{x} = (-1, 2)^\top$, $y = 2$. What is the optimal $\tilde{\mathbf{x}}^*$ that maximizes the objective in Eq. (OP)?

- ☐ $(-1 + \varepsilon, 2)^\top$
- ☐ $(-1 - \varepsilon, 2 - \varepsilon)^\top$
- ☐ $(-1 - \varepsilon, 2)^\top$
- ☐ $(-1 + \varepsilon, 2 + \varepsilon)^\top$
- ☒ Other

Solution: The optimal $\boldsymbol{\delta}^*$ is equal to $(-\varepsilon, \varepsilon)^\top$. This corresponds to $\tilde{\mathbf{x}}^* = \mathbf{x} + \boldsymbol{\delta}^* = (-1 - \varepsilon, 2 + \varepsilon)^\top$ which corresponds to the answer "Other".



KNN

Question 14 The KNN algorithm needs a notion of distance to assess which points are “nearest”. Identify the distance measures that can be used in the KNN algorithm

- (a) Euclidean Distance : distance associated to the L_2 norm $\|\mathbf{x}\|_2 := \sqrt{x_1^2 + \dots + x_D^2}$
- (b) Manhattan Distance : distance associated to the L_1 norm $\|\mathbf{x}\|_1 := |x_1| + \dots + |x_D|$
- (c) Distance associated to the L_4 norm $\|\mathbf{x}\|_4 := (|x_1|^4 + \dots + |x_D|^4)^{1/4}$

- ☐ only a and c
- ☐ only b
- ☐ only c
- ☒ a, b and c
- ☐ only b and c
- ☐ only a and b
- ☐ only a

Solution: The similarity measure is only an algorithmic choice that should be made in the KNN algorithm.

Linear Regression

Question 15 Assume we are doing linear regression with Mean-Squared Loss and L2-regularization on four one-dimensional data points. Our prediction model can be written as $f(x) = ax + b$ and the optimization problem can be written as

$$a^*, b^* = \underset{a, b}{\operatorname{argmin}} \sum_{n=1}^4 [y_n - f(x_n)]^2 + \lambda a^2.$$

Assume that our data points are $Y = [1, 3, 2, 4]$ and $X = [-2, -1, 0, 3]$. For example $y_1 = 1$ and $x_1 = -2$. What is the optimal value for the bias, b^* ?

- ☐ 2
- ☐ None of the above answers.
- ☐ Depends on the value of λ
- ☐ 3
- ☒ 2.5

Solution: If we take the derivative of the loss w.r.t b and set it to zero we get $\sum_{i=1}^4 [y_i - ax_i - b] = 0$. Since $\sum_{i=1}^4 ax_i = 0$, the optimal value for b is equal to the mean of the target values, $b = \frac{1+3+2+4}{4} = 2.5$.



Subgradients

Question 16 Consider the Parametric ReLU function defined as

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ ax & \text{otherwise} \end{cases}$$

where $a \in \mathbb{R}$ is an arbitrary number. Which of the following statements is true regarding the subgradients of $f(x)$ at $x = 0$?

- ☐ If a subgradient exists, then it is not unique.
- ☐ A subgradient does not exist at $x = 0$.
- ☐ A subgradient exists even though $f(x)$ is not necessarily differentiable at $x = 0$.
- ☒ None of the mentioned answers.

Solution: For values of $a > 1$ the function $f(x)$ is strictly concave and hence a subgradient does not exist. For $a = 1$ we have a unique subgradient. Hence the correct answer is None of the above answers.

Logistic regression

Consider a binary classification task as in Figure 1, which consists of 14 two-dimensional linearly separable samples (circles corresponds to label $y = 1$ and pluses corresponds to label $y = 0$). We would like to predict the label $y = 1$ of a sample (x_1, x_2) when the following holds true

$$\Pr(y = 1|x_1, x_2, w_1, w_2) = \frac{1}{1 + \exp(-w_1x_1 - w_2x_2)} > 0.5$$

where w_1 and w_2 are parameters of the model.




Figure 1: Two-dimensional dataset.



Question 17 If we obtain the (w_1, w_2) by optimizing the following objective

$$-\sum_{n=1}^N \log \Pr(y_n | x_{n1}, x_{n2}, w_1, w_2) + \frac{C}{2} w_2^2$$

where C is very large, then the decision boundary will be close to which of the following lines?

- ☐ $x_1 + x_2 = 0$
- ☒ $x_1 = 0$
- ☐ $x_2 = 0$
- ☐ $x_1 - x_2 = 0$

Solution: When C is very large, we are essentially optimizing over $\frac{C}{2} w_2^2$ and thus w_2 is close to 0. Then $\Pr(y = 1 | x_1, x_2, w_1, w_2) = \frac{1}{1 + \exp(-w_1 x_1)}$ indicates that the classifier outputs 1 or 0 depending on the sign of x . Therefore the decision boundary is $x_1 = 0$.

Recommender systems and word vectors

Question 18 What *alternates* in Alternating Least Squares for Matrix Factorization for a movie recommender system?

- ☐ recommendation steps and optimization steps
- ☐ updates based on different movie rating examples from the training set
- ☐ expectation steps and maximization steps
- ☒ updates to user embeddings and updates to movie embeddings

Question 19 Which NLP model architectures can differentiate between the sentences “I have to read this book.” and “I have this book to read.”?

- (a) a convolutional model based on word2vec vectors
 - (b) a recurrent neural network based on GloVe word vectors
 - (c) a bag-of-words model based on GloVe word vectors
- ☐ only b and c
 - ☐ only b
 - ☐ only a and c
 - ☐ a, b and c
 - ☐ only c
 - ☐ only a
 - ☒ only a and b



Generative Networks

Question 20 Consider a Generative Adversarial Network (GAN) which successfully produces images of goats. Which of the following statements is **false**?

- ☐ After the training, the discriminator loss should ideally reach a constant value.
- ☐ The generator aims to learn the distribution of goat images.
- ☐ The generator can produce unseen images of goats.
- ☒ The discriminator can be used to classify images as goat vs non-goat.

Solution: “*classify goat vs non-goat...*” is FALSE, because the discriminator classifies images into real or fake.

Clustering

Question 21 For any vector $\mathbf{v} \in \mathbb{R}^D$ let $\|\mathbf{v}\|_2 := \sqrt{v_1^2 + \dots + v_D^2}$ denote the Euclidean norm. Let $\mathbf{x}_1, \dots, \mathbf{x}_N \in \mathbb{R}^D$ be a dataset of $N \geq 2$ distinct points. For any integer $K \geq 1$, consider the following value:

$$L_K^* = \inf_{\substack{\mu_1, \dots, \mu_K \in \mathbb{R}^D \\ z_{nk} \in \{0,1\} \text{ s.t. } \sum_{k=1}^K z_{nk} = 1}} \frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K z_{nk} \|\mathbf{x}_n - \mu_k\|_2^2$$

The statements below are all true except one. Which of the following statement is *false*?

- ☐ For $2 \leq K < N$ and with the initial means randomly chosen as K data points, the K -means algorithm with K clusters is **not** guaranteed to reach the optimal L_K^* loss value.
- ☐ For $2 \leq K < N$, L_K^* is hard to compute.
- ☐ For $K \geq N$, $L_K^* = 0$.
- ☒ The sequence $(L_K^*)_{1 \leq K \leq N}$ is **not** necessarily strictly decreasing.
- ☐ L_1^* corresponds to the population variance of the dataset.

Solution: $(L_K^*)_{1 \leq K \leq N}$ is necessarily strictly decreasing. Since all the datapoints are distinct, we can always strictly improve the loss by assigning $\mu_{K+1} = \mathbf{x}_i$ where $x_i \neq \mu_k$ for all $1 \leq k \leq K$.



Question 22 You are given the data $(x_n, y_n)_{1 \leq n \leq N} \in \mathbb{R}^2$ illustrated in Figure 2 which you want to cluster into an inner ring and an outer ring (hence a number of clusters $K = 2$). Which of the following statement(s) is/are correct?

- (a) There exists some initialization such that K -means clustering succeeds.
- (b) There exists an appropriate feature expansion such that K -means (with standard initialization) succeeds.
- (c) There exists an appropriate feature expansion such that the Expectation Maximization algorithm (with standard initialization) for a Gaussian Mixture Model succeeds.

☐ Only a and c

☐ Only a

☐ All of them

☒ Only b and c

☐ None of them

☐ Only a and b

☐ Only c

☐ Only b

Solution: For example, $(x_i, y_i, c\sqrt{x_i^2 + y_i^2})_{1 \leq i \leq n}$ for c big enough will easily be clusterable using K -means or GMM.




Figure 2: Two-dimensional dataset



Linear Algebra

Question 23 Given a matrix \mathbf{X} of shape $D \times N$ with a singular value decomposition (SVD), $\mathbf{X} = \mathbf{U}\mathbf{S}\mathbf{V}^\top$, suppose \mathbf{X} has rank K and $\mathbf{A} = \mathbf{X}\mathbf{X}^\top$


Which one of the following statements is **false**?

- ☒ The eigenvalues of \mathbf{A} are the singular values of \mathbf{X}
- ☐ \mathbf{A} is positive semi-definite, i.e all eigenvalues of \mathbf{A} are non-negative
- ☐ The eigendecomposition of \mathbf{A} is also a singular value decomposition (SVD) of \mathbf{A}
- ☐ A vector \mathbf{x} that can be expressed as a linear combination of the last $D - K$ columns of \mathbf{U} , i.e $\mathbf{x} = \sum_{i=K+1}^D \mathbf{w}_i \mathbf{u}_i$ (where \mathbf{u}_i is the i -th column of \mathbf{U}), lies in the null space of \mathbf{X}^\top

Solution:

- “ \mathbf{A} is PSD...” and “The eigendecomposition of...” are true because $\mathbf{A} = \mathbf{U}\mathbf{S}^2\mathbf{U}^\top$.
- “The eigenvalues of... are the singular values...” is false because the eigenvalues of \mathbf{A} are the square of the singular values of \mathbf{X} .
- “A vector \mathbf{x} that...” is true.

Proof: $X^\top = VS^\top U^\top$, $X^\top \vec{x}_i = X^\top \sum_{i=K+1}^D w_i \vec{u}_i = \sum_{i=K+1}^D w_i (X^\top \vec{u}_i) = \sum_{i=K+1}^D w_i (VS^\top \vec{e}_i)$, the last equality follows from the fact that \vec{u}_i is orthogonal to every row of U^\top except the i th row i.e $U^\top \vec{u}_i = \vec{e}_i$. As X^\top has rank K , then S^\top has non-zeros along the diagonal only up to the k th row / column, i.e $S^\top \vec{e}_i = 0$ for $i \geq k + 1$. Therefore $X^\top \vec{x}_i = 0$



Second part: true/false questions

For each question, mark the box (without erasing) TRUE if the statement is **always true** and the box FALSE if it is **not always true** (i.e., it is sometimes false).

Question 24 (Neural Networks) Weight sharing allows CNNs to deal with image data without using too many parameters. However, weight sharing increases the variance of a model.

☐ TRUE ☒ FALSE

Solution: False, weight sharing increases bias.

Question 25 (Kernels) For any vector $\mathbf{v} \in \mathbb{R}^D$ let $\|\mathbf{v}\|_2 := \sqrt{v_1^2 + \dots + v_D^2}$ denote the Euclidean norm. Define the function $k(\mathbf{x}, \mathbf{x}') = \frac{1}{1 - \mathbf{x}^\top \mathbf{x}'}$, on the set $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^D$ such that $\|\mathbf{x}\|_2 < 1$ and $\|\mathbf{x}'\|_2 < 1$. The function $k(\mathbf{x}, \mathbf{x}')$ is a valid kernel.

☒ TRUE ☐ FALSE

Solution: True. It is a valid kernel, because $k(\mathbf{x}, \mathbf{x}') = \frac{1}{1 - \mathbf{x}^\top \mathbf{x}'} = \sum_{i=0}^{\infty} (\mathbf{x} \mathbf{x}')^i$

Question 26 (Bias/Variance Decomposition) Consider a linear regression model where the data is generated by input \mathbf{x} and output $y = \mathbf{w}^\top \mathbf{x} + \varepsilon$, where \mathbf{w} is a fixed vector and ε is a Gaussian noise with zero mean and σ^2 variance. Then there is no machine learning algorithm that can achieve a training error lower than σ^2 .

☐ TRUE ☒ FALSE

Solution: False. It is the true error that cannot be less than σ^2 instead of the training error. Moreover, overfitting is always an option, and it is often possible to have a small training error.

Question 27 (Logistic regression) Consider a binary classification task $L(\mathbf{w}) = \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}})$ with linearly separable samples and $y \in \{0, 1\}$. Then there exists an optimal \mathbf{w} with exact 0 loss and 100% training accuracy.

☐ TRUE ☒ FALSE

Solution: False. While there exists \mathbf{w} which has 100% training accuracy, exact 0 loss is not attainable.



Question 28 (Linear Models) For any vector $\mathbf{v} \in \mathbb{R}^D$ let $\|\mathbf{v}\|_2 := \sqrt{v_1^2 + \dots + v_D^2}$ denote the Euclidean norm. For $\mathbf{y} \in \mathbb{R}^D$, $\mathbf{X} \in \mathbb{R}^{D \times D}$, the solution of the least squares problem:

$$\mathbf{w}^* = \underset{\mathbf{w} \in \mathbb{R}^D}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2$$

is always unique.

☐ TRUE ☒ FALSE

Solution: False. This is true only if \mathbf{X} is full-rank. Suppose that

$$\mathbf{X} = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}, \text{ and } \mathbf{y} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

then $\mathbf{w}_1^* = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$ and $\mathbf{w}_2^* = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$ are both solutions.

Question 29 (PCA) Your friend performed Principal Component Analysis on some data and claimed that he retained at least 95% of the variance using k principal components. This is equivalent to $\frac{\sum_{i>k+1} \lambda_i}{\sum_i \lambda_i} \leq 0.05$, where $\lambda_1, \dots, \lambda_k, \dots$ are the eigenvalues associated to each principal component, sorted in a *non-increasing* order.

☒ TRUE ☐ FALSE

Solution: True. The variance explained with k principal components is 95%. This means that $\frac{\sum_{i \leq k} \lambda_i}{\sum_i \lambda_i}$ is at least 0.95. Hence, the claim is true.

Question 30 (Adversarial robustness) Let $\|\mathbf{v}\|_\infty := \max_i |v_i|$ denote the ℓ_∞ -norm. Assume a binary classification problem with $y \in \{-1, 1\}$. Then the adversarial zero-one loss $\max_{\delta: \|\delta\|_\infty \leq \epsilon} \mathbb{1}_{yf(\mathbf{x}+\delta) \leq 0}$ ($\mathbb{1}_C$ is equal to 1 if the condition C is true and to 0 if C is false) is always upper bounded by the adversarial hinge loss $\max_{\delta: \|\delta\|_\infty \leq \epsilon} \max\{0, 1 - yf(\mathbf{x} + \delta)\}$.

☒ TRUE ☐ FALSE

Solution: True. Since the exponential loss upper bounds the zero-one loss, taking the maximum over both of them preserves the ranking between them.

Question 31 (Expectation Maximization) If we run the Expectation Maximization (EM) algorithm on the log-likelihood for a Gaussian Mixture Model, then the sequence of log-likelihoods $\{\mathcal{L}(\boldsymbol{\theta}^{(t)})\}_{t \in \mathbb{N}}$ is guaranteed to be non-decreasing.

☒ TRUE ☐ FALSE

Solution: True. By definition of the algorithm, each parameter update increases the log-likelihood.



Question 32 (FastText Supervised Classifier) The FastText supervised classifier can be modelled as a two-layer linear (one hidden layer) neural network with a non-linear loss function at the output.

☒ TRUE ☐ FALSE

Solution: True. The FastText supervised classifier consists of two layers - embedding layer and linear classifier. The embedding layer can be thought of as another linear layer if the input is thought of as encoding word frequency.

Question 33 (Matrix Factorization) Consider the matrix factorization objective function $\frac{1}{2} \sum_{(d,n) \in \Omega} [x_{dn} - (\mathbf{W}\mathbf{Z}^\top)_{dn}]^2$. When minimizing this objective with SGD over the embedding matrices \mathbf{W} and \mathbf{Z} , you should initialize \mathbf{W} and \mathbf{Z} with zeros.

☐ TRUE ☒ FALSE

Solution: False. In fact, it will not work if you initialize the embeddings to zero. 'Zero' is a stationary point of the optimization problem. If you start all the parameters at zero, they will never change.

Question 34 (MSE and Neural Networks) The mean squared error (MSE) is convex w.r.t the parameters of a multi layer perceptron with more than one hidden layer and **linear** activation function.

☐ TRUE ☒ FALSE

Solution: False. Consider a simple two layer neural network with single neuron in each layer such that the output of the network is equal to $w_1 w_2 x$ where w_1 is the weight of the first layer and w_2 is the weight of the second layer. You can simply check that MSE loss is not convex w.r.t the parameters.



Third part, open questions

Answer in the space provided! Your answer must be justified with all steps. Leave the check-boxes empty, they are used for the grading.

PCA

Let $\mathbf{x}_1, \dots, \mathbf{x}_N$ be a dataset of N vectors in \mathbb{R}^D .

Question 35: (1 point.) What does it mean for the data vectors $\mathbf{x}_1, \dots, \mathbf{x}_N$ to be centered, as for principle component analysis (PCA) to be meaningful? Use the notation x_{nd} for individual entries.

☐ 0 ☒ 1

Solution: Data is centered, i.e. $\mathbb{E}[\mathbf{x}] = \mathbf{0}$ or in other words $\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \mathbf{0}$ or $\frac{1}{N} \sum_{n=1}^N x_{nd} = 0 \forall d$.

Question 36: (1 point.) Write down the covariance matrix of the dataset $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_N) \in \mathbb{R}^{D \times N}$, and state its dimensions. (Note that for PCA we assume data to be already centered, as in the previous question)

☐ 0 ☒ 1

Solution: $cov = \mathbf{X}\mathbf{X}^\top \in \mathbb{R}^{D \times D}$.

Now let \mathbf{x} be a random vector distributed according to the uniform distribution over the finite normalized dataset $\mathbf{x}_1, \dots, \mathbf{x}_N$ from above. Consider the problem of finding a unit vector, $\mathbf{w} \in \mathbb{R}^D$, such that the random variable $\mathbf{w}^\top \mathbf{x}$ has *maximal* variance.

Question 37: (1 point.) What is the variance of the random variable $\mathbf{w}^\top \mathbf{x}$ over the randomness of \mathbf{x} ?

☐ 0 ☒ 1

Solution: $\text{Var}[\mathbf{w}^\top \mathbf{x}] = \frac{1}{N} \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n)^2$

Question 38: (2 points.) Show that the solution of the problem of $\text{argmax}_{\mathbf{w}: \|\mathbf{w}\|=1} \text{Var}[\mathbf{w}^\top \mathbf{x}]$ is to set \mathbf{w} to be the first principle vector of $\mathbf{x}_1, \dots, \mathbf{x}_N$.


☐ 0 ☐ 1 ☒ 2

Solution: $\text{argmax}_{\mathbf{w}: \|\mathbf{w}\|=1} \text{Var}[\mathbf{w}^\top \mathbf{x}] = \frac{1}{N} \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n)^2 = \frac{1}{N} \sum_{n=1}^N \mathbf{w}^\top \mathbf{x}_n \mathbf{x}_n^\top \mathbf{w} = \frac{1}{N} \mathbf{w}^\top \mathbf{X}\mathbf{X}^\top \mathbf{w}$ is (by definition of Eigenvector) maximized if \mathbf{w} is the top eigenvector of $\mathbf{X}\mathbf{X}^\top$. (One can add some arguing why this gives the top singular vector of \mathbf{X})

Question 39: (1 point.) Explain in words what the above result says about how PCA relates to the dataset.

☐ 0 ☒ 1

Solution: The data has maximum variance when projected onto the first PC. (Mentioning 'max variance' alone is not sufficient, need to say what kind of variance, related to the data, direction).



The Perceptron

Setting

Let us consider a binary classification problem with a training set $S = \{(\mathbf{x}_n, y_n)\}_{n=1}^N$ such that:

$$\mathbf{x}_n \in \mathbb{R}^D, \text{ and } y_n \in \{-1, 1\}, \text{ for all } n = 1, \dots, N,$$

where N, D are integers such that $N, D \geq 1$.

We consider the Perceptron classifier which classifies $\mathbf{x} \in \mathbb{R}^D$ following the rule:

$$f_{\mathbf{w}, b}(\mathbf{x}) = \text{sgn}(\mathbf{w}^\top \mathbf{x} + b),$$

where $\mathbf{w} \in \mathbb{R}^D$ is the weight vector, $b \in \mathbb{R}$ is the threshold, and the sign function is defined as

$$\text{sgn}(z) = \begin{cases} +1 & \text{if } z \geq 0 \\ -1 & \text{if } z < 0 \end{cases}$$

Question 40: (1 point.) As seen in the course, explain how we can ignore the threshold b and only deal with classifiers passing through the origin, i.e., of the form $f_{\mathbf{w}}(\mathbf{x}) = \text{sgn}(\mathbf{w}^\top \mathbf{x})$.



Solution: We can view the threshold as an additional weight by adding the constant input 1 to the input \mathbf{x} . It amounts to consider the input $\tilde{\mathbf{x}}^\top = [\mathbf{x}^\top, 1]$ since $\tilde{\mathbf{x}}^\top [\mathbf{w}^\top, b] = \mathbf{x}^\top \mathbf{w} + b$.

For the remainder of the exercise, we proceed without this additive threshold, as explained in the previous question, and only consider classifiers of the form $f_{\mathbf{w}}(\mathbf{x}) = \text{sgn}(\mathbf{w}^\top \mathbf{x})$. We make the following two assumptions:

- **Bounded input:** There exists a real number $R \geq 0$ such that for all $n = 1, \dots, N$ we have

$$\|\mathbf{x}_n\|_2 \leq R,$$

where for any vector $\mathbf{z} \in \mathbb{R}^D$ we use $\|\mathbf{z}\|_2$ to refer to the Euclidean norm of \mathbf{z} , i.e., $\|\mathbf{z}\|_2 = \sqrt{\sum_{d=1}^D z_d^2}$.

- **Linearly separable data:** There exists $\mathbf{w}_* \in \mathbb{R}^D$ and a real number $\gamma > 0$ such that for all $n = 1, \dots, N$ we have

$$y_n \mathbf{w}_*^\top \mathbf{x}_n \geq \gamma.$$

We will use the Perceptron algorithm to train the Perceptron classifier and find a separating hyperplane. Let t denote the number of parameter updates we have performed and \mathbf{w}_t the weight vector after t updates. We initialize with $\mathbf{w}_0 = \mathbf{0}$. If this weight vector is already a separating hyperplane, we are done. If not, we pick an arbitrary point \mathbf{x}_i that is currently misclassified. This point is used to update the weight vector \mathbf{w}_t as

$$\mathbf{w}_{t+1} := \mathbf{w}_t + y_i \mathbf{x}_i \text{ where } y_i \mathbf{x}_i^\top \mathbf{w}_t \leq 0.$$

The algorithm is formally written below.

We will study the convergence of this learning algorithm.

Convergence

Question 41: (2 points.) For $t > 0$, show that when making the t^{th} update according to the Perceptron update, we have:

$$\mathbf{w}_*^\top \mathbf{w}_t \geq \mathbf{w}_*^\top \mathbf{w}_{t-1} + \gamma$$

**Algorithm 1:** Perceptron algorithm

```
 $t \leftarrow 0; \mathbf{w}_t \leftarrow 0;$ 
while there exists  $j \in \{1, \dots, N\}$  such that  $y_j \mathbf{x}_j^\top \mathbf{w}_t \leq 0$  do
    Pick an arbitrary  $i \in \{1, \dots, N\}$  such that  $y_i \mathbf{x}_i^\top \mathbf{w}_t \leq 0$  ;
     $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y_i \mathbf{x}_i$ ;
     $t \leftarrow t + 1$ 
end
return  $\mathbf{w}_t$ 
```

☐ 0 ☐ 1 ☒ 2

Solution: Let us assume that the t^{th} update is made on the sample i , we have then

$$\begin{aligned} \mathbf{w}_\star^\top \mathbf{w}_t &= \mathbf{w}_\star^\top (\mathbf{w}_{t-1} + y_i \mathbf{x}_i) \\ &= \mathbf{w}_\star^\top \mathbf{w}_{t-1} + y_i \mathbf{w}_\star^\top \mathbf{x}_i \\ &\geq \mathbf{w}_\star^\top \mathbf{w}_{t-1} + \gamma \end{aligned}$$

The first equation follows from the definition of the Perceptron update and the last equation follows from the assumption of linear separability .

Question 42: (1 point.) Show that it implies that

$$\mathbf{w}_\star^\top \mathbf{w}_t \geq t\gamma$$

☐ 0 ☒ 1

Solution: By induction on t we obtain $\mathbf{w}_\star^\top \mathbf{w}_t \geq \mathbf{w}_\star^\top \mathbf{w}_{t-1} + \gamma \geq \mathbf{w}_\star^\top \mathbf{w}_0 + t\gamma$, and by definition $\mathbf{w}_0 = 0$. Therefore $\mathbf{w}_\star^\top \mathbf{w}_t \geq t\gamma$.

Question 43: (1 point.) Using the previous question, derive a lower-bound on $\|\mathbf{w}_t\|_2$ depending on t , γ and $\|\mathbf{w}_\star\|_2$.

(We recall that a lower bound is a constant $C > 0$ such that $\|\mathbf{w}_t\|_2 \geq C$. Note that we will only accept non-trivial lower-bounds as correct answers.)

☐ 0 ☒ 1

Solution: Using the Cauchy-Schwarz inequality $\|\mathbf{w}_t\|_2 \|\mathbf{w}_\star\|_2 \geq \mathbf{w}_\star^\top \mathbf{w}_t \geq t\gamma$. Therefore $\|\mathbf{w}_t\|_2 \geq \frac{t\gamma}{\|\mathbf{w}_\star\|_2}$

Question 44: (3 points.) Derive the following upper bound on the squared norm $\|\mathbf{w}_t\|_2^2$:

$$\|\mathbf{w}_t\|_2^2 \leq \|\mathbf{w}_{t-1}\|_2^2 + R^2$$

☐ 0 ☐ 1 ☐ 2 ☒ 3

Solution: We still assume that the t^{th} is made on sample i , we have then

$$\begin{aligned} \|\mathbf{w}_t\|_2^2 &= \|\mathbf{w}_{t-1} + y_i \mathbf{x}_i\|_2^2 \\ &= \|\mathbf{w}_{t-1}\|_2^2 + 2y_i \mathbf{w}_{t-1}^\top \mathbf{x}_i + \|\mathbf{x}_i\|_2^2 \\ &\leq \|\mathbf{w}_{t-1}\|_2^2 + \|\mathbf{x}_i\|_2^2 \\ &\leq \|\mathbf{w}_{t-1}\|_2^2 + R^2 \end{aligned}$$



The first equation follows from the definition of the Perceptron update, the second equation follows from expanding the square, the third inequality follows from the fact that updates are made only on mistakes, i.e., $y_i \mathbf{w}_{t-1}^\top \mathbf{x}_i \leq 0$ whenever an update is made, and the last inequality follows from the boundedness assumption on the input.

Question 45: (1 point.) Show that it implies that

$$\|\mathbf{w}_t\|_2^2 \leq tR^2$$

☐ 0 ☒ 1

Solution: By induction on t we obtain $\|\mathbf{w}_t\|_2^2 \leq \|\mathbf{w}_{t-1}\|_2^2 + R^2 \leq \|\mathbf{w}_0\|_2^2 + tR^2$, and by definition $\mathbf{w}_0 = 0$. Therefore $\|\mathbf{w}_t\|_2^2 \leq tR^2$.

Question 46: (3 points.) Combine the results obtained in the previous questions to obtain an upper bound on t depending on the quantities R , γ and \mathbf{w}_* .

☐ 0 ☐ 1 ☐ 2 ☒ 3

Solution: Using Question 43 we get

$$\|\mathbf{w}_t\|_2 \geq \frac{t\gamma}{\|\mathbf{w}_*\|_2}.$$

From Question 45 we have

$$\|\mathbf{w}_t\|_2^2 \leq tR^2$$

Therefore combining both we obtain:

$$\frac{t^2\gamma^2}{\|\mathbf{w}_*\|_2^2} \leq \|\mathbf{w}_t\|_2^2 \leq tR^2$$

which yields

$$t \leq \frac{R^2\|\mathbf{w}_*\|_2^2}{\gamma^2}$$

Question 47: (2 points.) Qualitatively analyze the previous result. What have you proven? Interpret the dependency on R , γ and \mathbf{w}_* .

☐ 0 ☐ 1 ☒ 2

Solution: We have shown that the Perceptron algorithm terminates with a correct (separating) solution in finite time, and that this time is indirectly proportional to the margin γ^2 , and proportional to R^2 .

Perceptron and margin

Let us remind that we define the max-margin M_* as

$$M_* = \max_{\mathbf{w} \in \mathbb{R}^D, \|\mathbf{w}\|_2=1} M \text{ such that } y_n \mathbf{x}_n^\top \mathbf{w} \geq M \text{ for } n = 1, \dots, N$$

and a max-margin separating hyperplane $\bar{\mathbf{w}}$ as a solution of this problem:

$$\bar{\mathbf{w}} \in \arg \max_{\mathbf{w} \in \mathbb{R}^D, \|\mathbf{w}\|_2=1} M \text{ such that } y_n \mathbf{x}_n^\top \mathbf{w} \geq M \text{ for } i = 1, \dots, N$$



Question 48: (2 points.) Bound the number of perceptron updates t using the quantities R and M_* . Prove your result.

☐ 0 ☐ 1 ☒ 2

Solution: By definition of γ and M we have that $\gamma/\|\mathbf{w}_*\|_2 \leq M$ (1 point if proven properly). And therefore we obtain

$$t \leq \frac{R^2}{M^2}$$

Question 49: (1 point.) Does it imply that the output of the Perceptron algorithm is a max-margin separating hyperplane?

☐ 0 ☒ 1

Solution: No it does not, we have no clue to which separating hyperplane the Perceptron algorithm is converging to.

DRAFT